

Assignment 1 - Performance Analysis on a Quad-Core CPU

Problem 1 - Parallel Fractal Generation Using Threads

1. 修改 `mandelbrot.cpp` 中的代码以实现并行计算
2. 修改 `workerThreadStart` 的代码以实现
3. 研究针对不同图像的每个线程耗时，并解释原因
4. 实现一个8x的并行程序

Q&A

Is speedup linear in the number of cores used? In your writeup hypothesize why this is (or is not) the case?

我们针对 `view1` 进行测试

线程	耗时 /ms	加速比
1	74.321	1
2	37.048	2.01
3	24.824	2.99
4	18.942	3.92
5	16.844	4.41
6	14.119	5.26
7	12.197	6.093
8	10.648	6.979
9	17.736	4.19

这不是线性的，随着线程数增加，加速比增长速度逐渐下降。一直到9线程（本机CPU是8线程）时加速比降低。

按照理论，如果是线性关系，那么到达8线程的时候应该加速比接近8。考虑到我的本机CPU只有4核心，我不认为在5线程后每个线程性能是一致的。因此我们可以观察4线程时，可以注意到加速比为3.92，是比较接近4x加速的。

我个人认为，在5-8线程中，加速比增速下降的主要原因是cpu调度，cpu任务分配的耗时占了大头。

查阅了一些资料，是因为intel的超线程技术中，单线程性能不及一般线程，这也就解释了为什么5-8线程的时候加速比增速下滑了。

To confirm (or disprove) your hypothesis, measure the amount of time each thread requires to complete its work by inserting timing code at the beginning and end of `workerThreadStart()`. How do your measurements explain the speedup graph you previously created?

在这个问题中，我给每个核心都开了性能模式，因此速度会快一点

```
> ./mandelbrot --view 1 --threads 4
[mandelbrot serial]:          [71.300] ms
Wrote image file mandelbrot-v1-serial.ppm
[thread 3]: [18.512] ms
[thread 2]: [18.527] ms
[thread 1]: [18.593] ms
[thread 0]: [18.725] ms
[thread 2]: [18.523] ms
[thread 1]: [18.584] ms
[thread 3]: [18.573] ms
[thread 0]: [18.591] ms
[thread 3]: [18.553] ms
[thread 2]: [18.577] ms
[thread 1]: [18.655] ms
[thread 0]: [18.670] ms
[mandelbrot thread]:          [18.669] ms
Wrote image file mandelbrot-v1-thread-4.ppm
+++++                          (3.82x speedup from 4 threads)

> ./mandelbrot --view 2 --threads 4
[mandelbrot serial]:          [42.205] ms
Wrote image file mandelbrot-v2-serial.ppm
[thread 3]: [11.019] ms
[thread 2]: [11.066] ms
[thread 0]: [11.096] ms
[thread 1]: [11.224] ms
[thread 2]: [11.047] ms
[thread 3]: [11.029] ms
[thread 1]: [11.075] ms
[thread 0]: [11.124] ms
[thread 1]: [11.076] ms
[thread 2]: [11.236] ms
[thread 3]: [12.659] ms
[thread 0]: [12.844] ms
[mandelbrot thread]:          [11.173] ms
Wrote image file mandelbrot-v2-thread-4.ppm
+++++                          (3.78x speedup from 4 threads)
```

其实不明白为何验证非线性结论要求我计算每个线程的耗时

但是我发现了一个新问题：我开了4 threads,但是无论怎么改都有12个输出

后面改了不同的线程，1threads有3个输出，2就有6个，3就有9个，神奇啊（大雾）

Modify the mapping of work to threads to improve speedup to 8× on view 0 and almost 8× on views 1 and 2. In your writeup, describe your approach and report the final 16-thread speedup obtained. Also comment on the difference in scaling behavior from 4 to 8 threads versus 8 to 16 threads.

就是让我们改程序，以让它可以实现8x速度，但是cmu的高端电脑似乎是8核心16线程的，我的老破小是4核心8线程的，应该只要实现4x速度吧（大雾）

但是似乎已经改不到更快了，3.92x已经是极限了（悲）

Code

```
// Line 109-128
void *workerThreadStart(void *threadArgs) {
    WorkerArgs *args = static_cast<WorkerArgs *>(threadArgs);

    // TODO: Implement worker thread here.
    int threadNumRows = args->height / args->numThreads;
    if(args->height % args->numThreads != 0) {
        threadNumRows++;
    }
    // int threadStartRow = args->threadId * threadNumRows;
    // int threadTotalRows = threadNumRows;

    // mandelbrotSerial(args->x0, args->y0, args->x1, args->y1, args->width, args->
    >height, threadStartRow, threadTotalRows, args->maxIterations, args->output);

    for(int i = 0; i < threadNumRows; i++) {
        int threadStartRow = i * args->numThreads + args->threadId;
        if(threadStartRow < args->height) {
            mandelbrotSerial (args->x0, args->y0, args->x1, args->y1, args->width, args->
            >height, threadStartRow, 1, args->maxIterations, args->output);
        }
    }
    return NULL;
}

// Line 147 - 159
for (int i = 0; i < numThreads; i++) {
    args[i].threadId = i;
    // TODO: Set thread arguments here
    args[i].x0 = x0, args[i].y0 = y0;
    args[i].x1 = x1, args[i].y1 = y1;
    args[i].width = width;
    args[i].height = height;
    args[i].maxIterations = maxIterations;
    args[i].output = output;
    args[i].numThreads = numThreads;
}
```

Problem 2 - Vectorizing Code Using SIMD Intrinsics

- 1. 修改 `functions.cpp` 中的 `clampedExpVector()` 以实现并行快速幂
- 2. 修改 `functions.cpp` 中的 `arraySumVector()` 以实现并行数组求和

Q&A

How much does the vector utilization change as W changes? Explain the reason for these changes and the degree of sensitivity the utilization has on the vector width. Explain how the total number of vector instructions varies with W .

我们运行参数为 `./vrun -s 10000`

Vector Width	Vector Utilization	Total Vector Instructions
2	92.980082%	383929
4	92.812667%	196495
8	92.777550%	98739
16	92.776768%	49378
32	92.778057%	24651

我们可以发现的，在误差范围内，当向量宽度逐渐变大时，我们向量利用率是逐步下降的，但是下降速率逐渐降低
利用率下降的主要原因应该是，代码中我对余数部分的数组线性处理，而当宽度变大时，线性处理的部分也会增加

Code

我一开始是直接上面的快速幂改成向量的，只处理了除不尽的情况，然后当s 为奇数的时候会Fail
后面忘记看了什么资料发现别人的代码中是移走if(result > 4.18f)的代码到while结束，手动改了一下就过了
2023-03-31：重新回顾了一下这个代码，发现我们不需要线性处理，我们只需要设置一下掩码即可，修改后效率可以达到 `94.207317%` ，提升2%
下面这份是修改后的代码，修改前的话则是对余数部分线性处理

```
void clampedExpVector(float *values, int *exponents, float *output, int N) {
    // TODO: Implement your vectorized version of clampedExpSerial here
    __cmu418_vec_float x, result, xpower;
    __cmu418_vec_int y, tmp_y; // tmp_y stores y & 0x1

    __cmu418_vec_float exp = __cmu418_vset_float(4.18f);
    __cmu418_vec_int zero = __cmu418_vset_int(0);
    __cmu418_vec_int one = __cmu418_vset_int(1);

    __cmu418_mask maskAll, maskIsNegative, maskIsNotNegative, maskResultIsClamped;

    int cnt = N % VECTOR_WIDTH;
```

```

// if(cnt) clampedExpSerial(values, exponents, output, cnt); // Serial Solve

// Vector Solve
for (int i = 0; i < N; i += VECTOR_WIDTH) {
    maskAll = (i == 0 && cnt) ? _cmu418_init_ones(cnt) : _cmu418_init_ones();
    maskIsNegative = _cmu418_init_ones(0);

    _cmu418_vload_float(x, values + i, maskAll); // x = values[i];
    _cmu418_vset_float(result, 1.f, maskAll); // result = 1.f;
    _cmu418_vload_int(y, exponents + i, maskAll); // y =
exponents[i];
    _cmu418_vmove_float(xpower, x, maskAll); // xpower = x;

    while(_cmu418_vgt_int(maskIsNegative, y, zero, maskAll),
_cm418_cntbits(maskIsNegative)) // while(y > 0)
    {
        _cmu418_vbitand_int(tmp_y, y, one, maskAll); // tmp_y
= y & 0x1
        _cmu418_vgt_int(maskIsNegative, tmp_y, zero, maskAll); //
if(tmp_y) {
            _cmu418_vmult_float(result, result, xpower, maskIsNegative); //
result *= xpower
            // }
            _cmu418_vmult_float(xpower, xpower, xpower, maskAll); // xpower
*= xpower
            _cmu418_vshiftright_int(y, y, one, maskAll); // y >>=
1
        }

        _cmu418_vgt_float(maskResultIsClamped, result, exp, maskAll); // if(result >
4.18f) {
            _cmu418_vmove_float(result, exp, maskResultIsClamped); // result =
4.18f
            // }
            _cmu418_vstore_float(output + i, result, maskAll); // output[i] =
result

            if(i == 0 && cnt) i -= VECTOR_WIDTH;
        }
    }
}

```

这里有个未解之谜，关于Bonus的

```

float arraySumVector(float *values, int N) {
    // TODO: Implement your vectorized version here

    __cmu418_vec_float x, result;
    __cmu418_vec_float zero = _cmu418_vset_float(0.f);
    __cmu418_mask maskAll, maskIsNegative, maskIsNotNegative;

```

```

float sum[VECTOR_WIDTH];
float ans = 0.f;

for(int i = 0; i < VECTOR_WIDTH; i++)
{
    sum[i] = 0.f;
    // printf("sum[%d] = %f\n", i, sum[i]);
}

for(int i = 0; i < N; i += VECTOR_WIDTH)
{
    maskAll = _cmu418_init_ones();
    maskIsNegative = _cmu418_init_ones(0);

    _cmu418_vload_float(x, values + i, maskAll);           // x =
values[i];
    _cmu418_vadd_float(result, result, x, maskAll);       // result +=
x;
    _cmu418_vstore_float(sum + (i % VECTOR_WIDTH), result, maskAll); // sum[i] =
result;
}

for(int i = 0; i < VECTOR_WIDTH; i++) printf("sum[%d] = %f\n", i, sum[i]), ans +=
sum[i];

return ans;
}

```

这是我的这份代码，可以注意到我有一个printf的输出，当我注释掉输出时，会输出如下

```

ARRAY SUM (bonus)
sum[0] = -1.010000
sum[1] = -1.007556
sum[2] = -1.005328
sum[3] = -1.000470
sum[4] = -1.006793
sum[5] = -1.003835
sum[6] = -256138150223410355193484371034112.000000
sum[7] = -1.000535
Expected -8.032826, got -256138150223410355193484371034112.000000
.@@@ Failed!!!

```

当我没有注释输出时，会输出如下

```

ARRAY SUM (bonus)
sum[0] = 0.000000
sum[1] = 0.000000
sum[2] = 0.000000
sum[3] = 0.000000
sum[4] = 0.000000

```

```

sum[5] = 0.000000
sum[6] = 0.000000
sum[7] = 0.000000
sum[0] = -1.000000
sum[1] = -1.007556
sum[2] = -1.015328
sum[3] = -1.000470
sum[4] = -1.006793
sum[5] = -1.003835
sum[6] = -1.008310
sum[7] = -1.000535
Expected -8.032826, got -8.042827
.@@@ Failed!!!

```

??? 多一个printf会有这么大的影响嘛，不过最终答案的确也是错的，不知道发生了什么。太玄学了！

看了一下 `CMU418intrin.h`，我们肯定是要用上 `_cmu418_hadd_float()` 和 `_cmu418_interleave_float()` 的，因为这两个很突兀。第一个是实现 `[0 1 2 3] → [0+1 0+1 2+3 2+3]`，不好描述，就是按pair相加，第二个是实现把一个数组的奇数下标移到前半部分，偶数下标移动到后半部分

很明显嘛，就是循环使用这两个函数，类似二分，每次都会折半，题目中说了 `VECTOR_WIDTH` 是2的倍数，那么我们可以按位处理。

然后同时还发现了这行代码 `template <typename T> struct __cmu418_vec { T value[VECTOR_WIDTH]; };`，原来可以直接输出向量的某一个下标

那就好办了，不需要sum了

```

float arraySumVector(float *values, int N) {
    // TODO: Implement your vectorized version here

    __cmu418_vec_float x, result = _cmu418_vset_float(0.f);
    __cmu418_mask maskAll = _cmu418_init_ones();
    int cnt = VECTOR_WIDTH;

    for(int i = 0; i < N; i += VECTOR_WIDTH)
    {
        _cmu418_vload_float(x, values + i, maskAll);           // x =
values[i];
        _cmu418_vadd_float(result, result, x, maskAll);       // result +=
x;
    }

    while(cnt >= 1)
    {
        _cmu418_hadd_float(result, result);
        _cmu418_interleave_float(result, result);
    }

    return result.value[0];
}

```

最终结果

```
> ./vrun -s 8
CLAMPED EXPONENT (required)
Results matched with answer!
***** Printing Vector Unit Statistics *****
Vector Width:           8
Total Vector Instructions: 82
Vector Utilization:     93.140244%
Utilized Vector Lanes:  611
Total Vector Lanes:     656
***** Result Verification *****
Passed!!!

ARRAY SUM (bonus)
Passed!!!
```

Problem 3 - Using Instruction-level Parallelism in Fractal Generation

1. 理解并注释 `mandelbrot_ref_loop.s` 中的代码

Q&A

An annotated version of the assembly code for the main loop of mandel ref

我们可以在这里查询汇编指令: [x86 and amd64 instruction reference](#)

我们可以在这里查询一些其他资料: [并发原理 - CPU原子性指令 \(一\)](#)

那么其实就是查询汇编的答案了, 但是还是有点不会。

要知道一些基本信息, 例如x86-64的16个64位寄存器

Name	Description
%xmmi	128位寄存器, i为编号
%rip	指向下一条要执行的指令
%eax	accumulator, 加乘法指令缺省寄存器
%edi	destination indexm, 目标索引寄存器

An analysis of the latency bound of mandel ref and how the measured performance compares.

主要是提交一份Mandel Ref的延迟分析

这份代码中我觉得比较多的是xmm寄存器的操作, 查看Assignments的介绍中有提及: U0可以执行加乘除、U1可以执行加乘、U5可以做浮点加法, 但是只能使用x87浮点指令。U0/U1有4个时钟周期, U5有3个时钟周期。

但是我没有在[The microarchitecture of Intel and AMD CPUs](#)中对应页面找到解释（？）难道是更新了，似乎是在P155

这部分很薄弱，我的分析可能存在问题：

整个loop中涉及到了4次寄存器乘法，5次寄存器加法，还有2次jmp，latency至少是 $4 * 4 + 5 * 3 + 1 * 2 = 26$ 时钟周期

Considering the floating-point operations, what is the highest throughput bound imposed by the functional units?

考虑浮点运算，功能单元所施加的最高吞吐量是多少

这个超纲了啊!!!

Code

以下给出这份代码的注释版本，但是我感觉没有写好

```
# This is the inner loop of mandel_ref
# Parameters are passed to the function as follows:
#   %xmm0: c_re
#   %xmm1: c_im
#   %edi:  count
# Before entering the loop, the function sets registers
# to initialize local variables:
#   %xmm2: z_re = c_re
#   %xmm3: z_im = c_im
#   %eax:  i = 0
.L123:
    vmulss %xmm2, %xmm2, %xmm4      # xmm2 *= xmm4
    vmulss %xmm3, %xmm3, %xmm5      # xmm3 *= xmm5
    vaddss %xmm5, %xmm4, %xmm6      # xmm5 = xmm4 * xmm6
    vucomiss    .LC0(%rip), %xmm6    # Compare xmm1 and xmm6 and set the EFLAGS
flags accordingly.
    ja      .L126                    # if larger, jump to .L126
    vaddss %xmm2, %xmm2, %xmm2      # xmm2 += xmm2
    addl    $1, %eax                # count++
    cmpl    %edi, %eax              # Set condition codes for jne below
    vmulss %xmm3, %xmm2, %xmm3      # xmm3 *= xmm2
    vsubss %xmm5, %xmm4, %xmm2      # xmm5 = xmm4 - xmm2
    vaddss %xmm3, %xmm1, %xmm3      # xmm3 += xmm1
    vaddss %xmm2, %xmm0, %xmm2      # xmm2 += xmm0
    jne     .L123
```

Problem 4 - Parallel Fractal Generation Using ISPC

出了点小插曲，我的电脑中因为更新已经删除了 `sys/sysctl.h`，不过似乎注释掉 `common/tasksys.h` 中的include就可以编译了，那导入这个库干嘛..

Problem 4, Part 1. A Few ISPC Basics

1. 编译程序，并提出自己认为的理论值

Q&A

What is the maximum speedup you expect given what you know about these CPUs? Why might the number you observe be less than this ideal?

首先，既然是 8-wide AVX vector instructions，那理论加速比应当为8x，但是实际上我跑的理论加速比只有4.48x

我认为没有达到实际加速比是因为调度问题，毕竟不可能直接到8x，但是4.48x的差距还是太大了，说明不止是调度问题

查阅了网上的标准答案 他人的想法，负载不均衡的问题可能是比较正确的，我们实际上跑的这个图不是均匀的，那么就会造成有的快有的慢。

Problem 4, Part 2. Combining instruction-level and SIMD parallelism

Q&A

How much speedup does this two-way parallelism give over the regular ISPC version? Does it vary across different inputs (i.e., different `--views`)? When is it worth the effort?

这道题做崩了，反向加速了，但是可以注意到在视图不同的时候，反向加速效果也是不一样的，但是我感觉我的代码没有很大的问题，实在是不知道为啥反向加速了，难道是自己电脑CPU太古早了？

```
> ./mandelbrot_ispc -v 1
[mandelbrot serial]:          [183.684] ms
Wrote image file mandelbrot-1-serial.ppm
[mandelbrot ispc]:           [44.045] ms
Wrote image file mandelbrot-1-ispc.ppm
[mandelbrot ispc parallel]:    [66.698] ms
Wrote image file mandelbrot-1-ispc-par.ppm
                                (4.17x speedup from ISPC)
                                (2.75x speedup from ISPC+parallelism)

> ./mandelbrot_ispc -v 2
[mandelbrot serial]:          [107.956] ms
Wrote image file mandelbrot-2-serial.ppm
[mandelbrot ispc]:           [30.914] ms
Wrote image file mandelbrot-2-ispc.ppm
[mandelbrot ispc parallel]:    [43.129] ms
Wrote image file mandelbrot-2-ispc-par.ppm
```

```

(3.49x speedup from ISPC)
(2.50x speedup from ISPC+parallelism)
> ./mandelbrot_ispc -v 3
[mandelbrot serial]:          [260.036] ms
Wrote image file mandelbrot-3-serial.ppm
[mandelbrot ispc]:           [65.158] ms
Wrote image file mandelbrot-3-ispc.ppm
[mandelbrot ispc parallel]:   [89.083] ms
Wrote image file mandelbrot-3-ispc-par.ppm
(3.99x speedup from ISPC)
(2.92x speedup from ISPC+parallelism)
> ./mandelbrot_ispc -v 4
[mandelbrot serial]:          [257.352] ms
Wrote image file mandelbrot-4-serial.ppm
[mandelbrot ispc]:           [64.656] ms
Wrote image file mandelbrot-4-ispc.ppm
[mandelbrot ispc parallel]:   [88.640] ms
Wrote image file mandelbrot-4-ispc-par.ppm
(3.98x speedup from ISPC)
(2.90x speedup from ISPC+parallelism)

```

Code

```

export void mandelbrot_ispc_par2(uniform float x0, uniform float y0,
                                uniform float x1, uniform float y1,
                                uniform int width, uniform int height,
                                uniform int maxIterations,
                                uniform int output[]) {
    uniform float dx = (x1 - x0) / width;
    uniform float dy = (y1 - y0) / height;

    // TODO: Write ISPC code that will use function mandel_par2 to process
    // two rows on each pass.
    // You should use the foreach construct.
    // You should handle the case where the height is not a multiple
    // of 2.

    uniform int numRowsPerPass = (height + 1) / 2;

    foreach (i = 0 ... numRowsPerPass, j = 0 ... width) {
        int row0 = i * 2;
        int row1 = min(row0 + 1, height - 1);

        float c_re0 = x0 + j * dx;
        float c_im0 = y0 + row0 * dy;
        float c_re1 = x0 + j * dx;
        float c_im1 = y0 + row1 * dy;
    }
}

```

```

        mandel_par2(c_re0, c_im0, c_re1, c_im1, maxIterations, output + row0 * width + j,
output + row1 * width + j);
    }
}

```

Problem 4, Part 3: ISPC Tasks

1. 尝试不同的BLOCK大小与TASK数量，以超过20x加速

最终是 20.99x 加速，参数是

```

#define BLOCK_WIDTH 400
#define BLOCK_HEIGHT 200
uniform int taskCount = 16;

```

不过这个20.99x也蛮极限的，可能原因是CPU是4核心8线程的，16task其实是有点高了

有一个问题是说

Problem 5 - Iterative Cubic Root

1. 运行迭代立方根程序，看结果
2. 修改 `data.cpp` 中的内容，调优

```

> ./cuberoot
[cuberoot serial]:          [3658.882] ms
[cuberoot ispc]:           [884.804] ms
[cuberoot task ispc]:      [134.455] ms
                             (4.14x speedup from ISPC)
                             (27.21x speedup from task ISPC)

```

Q&A

Modify the function `initGood/initBad()` in the file `data.cpp` to generate data that will yield a very high relative speedup of the ISPC implementations.

Does your modification improve SIMD speedup? Does it improve multi-core speedup? Please explain why.

我最多可以跑到40x左右的加速，这时候 `value[i] = 1.9999999`，似乎可以无限逼近2

```

> ./cuberoot -d g
[cuberoot serial]:          [7345.921] ms
[cuberoot ispc]:           [1190.029] ms
[cuberoot task ispc]:      [178.646] ms
                             (6.17x speedup from ISPC)
                             (41.12x speedup from task ISPC)

```

最慢可以达到2.3x的加速，这时候 `value[i] = 1.0`

```
[cuberoot serial]:          [80.059] ms
[cuberoot ispc]:           [43.914] ms
[cuberoot task ispc]:      [34.633] ms
                             (1.82x speedup from ISPC)
                             (2.31x speedup from task ISPC)
```

实际上修改为1就可以实现最慢了，-1应该也可以，可以观察给的图形，在正负1.0处迭代最少，在无限逼近2的时候迭代次数徒增

要解释的话，我们可以说明ispc可以更好的处理数学计算？加速数学计算的过程吧。

这个问题5似乎比前面几个来的简单多啊。。这个看起来得放在第一问（）